

# Preparing for Future Pandemics Vaccine Needs

This project aims to help health providers prepare for the vaccination needs for future pandemics. From community clinics to health care providers with hospitals and offices across regions, they will need to know the amount of vaccines needed. The goals are to understand the indicators for receiving pandemic vaccines and build a model that will predict those who will get a pandemic vaccine.

- [Business Understanding](#)
- [Data Understanding](#)
- [Data Preperation](#)
- [Modeling](#)
- [Evaluation](#)

## Business Question

How many vaccines healthcare providers need to purchase or request, depending on the national distribution plan, so that all patients who want a vaccine receive one and will not collect a large surplus.

To answer this question we are going to build a model trained on data of who received the h1n1 vaccine, to use to predict individuals who will likely receive a future pandemic vaccination.

While overall accuracy is important, because health care providers would rather have slightly more vaccines than needed, rather than being short and not having vaccines for individuals that requested them, we will focus on achieving a high recall score.

## Data Understanding

The data comes from an over 26,000 person phone survey conducted in 2010, a year after the H1N1 outbreak, in which participants were asked about receiving the H1N1 vaccine, the seasonal flu vaccine, opinions about vaccines, behaviors around transmitting illness, and demographic information.

First join the training data and the targets and take a look at the dataset.

```
In [1]: 1 # Import Libraries
2
3 import pandas as pd
4 import numpy as np
5 from matplotlib import pyplot as plt
6 import seaborn as sns
7
8 from sklearn.model_selection import train_test_split, GridSearchCV, cross_val_score
9 from sklearn.pipeline import Pipeline
10 from sklearn.preprocessing import StandardScaler, OneHotEncoder, FunctionTransformer, OrdinalEncoder
11 from sklearn.impute import SimpleImputer
12 from sklearn.compose import ColumnTransformer, make_column_selector as selector
13 from sklearn.linear_model import LogisticRegression
14 from sklearn.ensemble import RandomForestClassifier, GradientBoostingClassifier
15 from sklearn.neighbors import KNeighborsClassifier
16 from sklearn.naive_bayes import MultinomialNB, GaussianNB
17
18 from sklearn.metrics import plot_confusion_matrix, recall_score, \
19     accuracy_score, precision_score, f1_score
20
21 from imblearn.over_sampling import SMOTE
22 from imblearn.pipeline import Pipeline as ImPipeline
23
24 from sklearn.dummy import DummyClassifier
```

```
In [2]: 1 # Set Options
2
3 pd.set_option('max_columns', None)
4 sns.set_palette("colorblind")
5 sns.set_style("darkgrid")
6
```

```
In [3]: 1 # Read in csv data files
        2
        3 df = pd.read_csv('data/training_set_features.csv')
        4 df_tars = pd.read_csv('data/training_set_labels.csv')
        5
        6 df_tars.head()
```

Out[3]:

	respondent_id	h1n1_vaccine	seasonal_vaccine
0	0	0	0
1	1	0	1
2	2	0	0
3	3	0	1
4	4	0	0

```
In [4]: 1 # Join together the files on "respondent_id"
        2
        3 df = df.join(df_tars, on='respondent_id', rsuffix='_tars')
```

```
In [5]: 1 # Sanity check ids match
        2
        3 df.loc[:, ['respondent_id', 'respondent_id_tars']]
```

Out[5]:

	respondent_id	respondent_id_tars
0	0	0
1	1	1
2	2	2
3	3	3
4	4	4
...	...	...
26702	26702	26702
26703	26703	26703
26704	26704	26704
26705	26705	26705
26706	26706	26706

26707 rows × 2 columns

```
In [6]: 1 # Drop repeated column and check shape, then investigate the dataframe
        2
        3 df = df.drop('respondent_id_tars', axis=1)
        4 df.shape
```

Out[6]: (26707, 38)

```
In [7]: 1 df.head()
```

Out[7]:

	h1n1_risk	opinion_h1n1_sick_from_vacc	opinion_seas_vacc_effective	opinion_seas_risk	opinion_seas_sick_from_vacc	age_group	education	race	sex
	1.0	2.0	2.0	1.0	2.0	55 - 64 Years	< 12 Years	White	Female
	4.0	4.0	4.0	2.0	4.0	35 - 44 Years	12 Years	White	Male
	1.0	1.0	4.0	1.0	2.0	18 - 34 Years	College Graduate	White	Male
	3.0	5.0	5.0	4.0	1.0	65+ Years	12 Years	White	Female
	3.0	2.0	3.0	1.0	4.0	45 - 54 Years	Some College	White	Female

In [8]:

```
1 df.info()

<class 'pandas.core.frame.DataFrame'>
RangeIndex: 26707 entries, 0 to 26706
Data columns (total 38 columns):
 #   Column                                Non-Null Count  Dtype
---  -
 0   respondent_id                        26707 non-null  int64
 1   h1n1_concern                        26615 non-null  float64
 2   h1n1_knowledge                      26591 non-null  float64
 3   behavioral_antiviral_meds           26636 non-null  float64
 4   behavioral_avoidance                 26499 non-null  float64
 5   behavioral_face_mask                26688 non-null  float64
 6   behavioral_wash_hands               26665 non-null  float64
 7   behavioral_large_gatherings         26620 non-null  float64
 8   behavioral_outside_home             26625 non-null  float64
 9   behavioral_touch_face               26579 non-null  float64
10  doctor_recc_h1n1                   24547 non-null  float64
11  doctor_recc_seasonal                24547 non-null  float64
12  chronic_med_condition               25736 non-null  float64
13  child_under_6_months               25887 non-null  float64
14  health_worker                      25903 non-null  float64
15  health_insurance                   14433 non-null  float64
16  opinion_h1n1_vacc_effective          26316 non-null  float64
17  opinion_h1n1_risk                    26319 non-null  float64
18  opinion_h1n1_sick_from_vacc          26312 non-null  float64
19  opinion_seas_vacc_effective          26245 non-null  float64
20  opinion_seas_risk                    26193 non-null  float64
21  opinion_seas_sick_from_vacc          26170 non-null  float64
22  age_group                           26707 non-null  object
23  education                           25300 non-null  object
24  race                                26707 non-null  object
25  sex                                 26707 non-null  object
26  income_poverty                     22284 non-null  object
27  marital_status                     25299 non-null  object
28  rent_or_own                        24665 non-null  object
29  employment_status                  25244 non-null  object
30  hhs_geo_region                     26707 non-null  object
31  census_msa                         26707 non-null  object
32  household_adults                   26458 non-null  float64
33  household_children                 26458 non-null  float64
34  employment_industry                13377 non-null  object
35  employment_occupation              13237 non-null  object
36  h1n1_vaccine                       26707 non-null  int64
37  seasonal_vaccine                   26707 non-null  int64
dtypes: float64(23), int64(3), object(12)
memory usage: 7.7+ MB
```

In [9]:

```
1 df.describe()
```

Out[9]:

m_vacc	opinion_seas_vacc_effective	opinion_seas_risk	opinion_seas_sick_from_vacc	household_adults	household_children	h1n1_vaccine	seasonal_vaccine
.000000	26245.000000	26193.000000	26170.000000	26458.000000	26458.000000	26707.000000	26707.000000
.357670	4.025986	2.719162	2.118112	0.886499	0.534583	0.212454	0.465608
.362766	1.086565	1.385055	1.332950	0.753422	0.928173	0.409052	0.498825
.000000	1.000000	1.000000	1.000000	0.000000	0.000000	0.000000	0.000000
.000000	4.000000	2.000000	1.000000	0.000000	0.000000	0.000000	0.000000
.000000	4.000000	2.000000	2.000000	1.000000	0.000000	0.000000	0.000000
.000000	5.000000	4.000000	4.000000	1.000000	1.000000	0.000000	1.000000
.000000	5.000000	5.000000	5.000000	3.000000	3.000000	1.000000	1.000000

In [10]:

```
1 df.select_dtypes(include='object').describe()
```

Out[10]:

	age_group	education	race	sex	income_poverty	marital_status	rent_or_own	employment_status	hhs_geo_region	census_msa	employment
count	26707	25300	26707	26707	22284	25299	24665	25244	26707	26707	
unique	5	4	4	2	3	2	2	3	10	3	
top	65+ Years	College Graduate	White	Female	<= \$75,000, Above Poverty	Married	Own	Employed	Izgxpyit	MSA, Not Principle City	
freq	6843	10097	21222	15858	12777	13555	18736	13560	4297	11645	

The dataset now has 26707 rows (survey respondents), and 38 columns (variables including id and targets).

It has 12 objects and 26 numeric indicators. Indicators related to behavioral questions are binary, indicators related to opinion questions are a five point scale, numerical and string indicators related to demographics are of varying numbers of response choices, and respondent ID is a unique identifier. The targets are binary, 0 or 1.

'health\_insurance', 'employment\_industry', and 'employment\_occupation' all have over 50% nulls. Other columns contain a small percentage of nulls.

## Data Preperation


What are the numbers for the targets? How many respondents received the vaccines?

```
In [11]: 1 # Look at the percent of respondents who have recieved the H1N1 Vaccine
2
3 # Get Proportions
4 vax_prop = df.h1n1_vaccine.value_counts(normalize=True)
5 print(vax_prop)
6
7
8 # Chart
9 fig, ax = plt.subplots()
10
11 sns.barplot(x=vax_prop.index, y=vax_prop.values*100)
12
13 ax.set_xlabel('')
14 ax.set_ylabel('Percent Vaccine')
15 ax.set_title("H1N1 Vaccination")
16
17 ax.set_xticklabels(['No H1N1 Vax', 'H1N1 Vax'])
18
19 plt.savefig('images/vaccine_percents.png', bbox_inches='tight', dpi=300)
```

```
0    0.787546
```


```
1    0.212454
```

```
Name: h1n1_vaccine, dtype: float64
```



```
In [12]: 1 # Look at the proportion of respondents who have recieved seasonal flu Vaccine for comparison
2
3 # Get Proportions
4 vax_prop_seasonal = df.seasonal_vaccine.value_counts(normalize=True)
5 print(vax_prop_seasonal)
6
7
8 # Chart
9 fig, ax = plt.subplots()
10
11 sns.barplot(x=vax_prop_seasonal.index, y=vax_prop_seasonal.values*100)
12
13 ax.set_xlabel('')
14 ax.set_ylabel('Percent Vaccine')
15 ax.set_title("Seasonal Flu Vaccination")
16
17 ax.set_xticklabels(['No Seasonal Vax', 'Seasonal Vax'])
18
19 plt.savefig('images/seasonal_vax_percents.png', bbox_inches='tight', dpi=300)
```

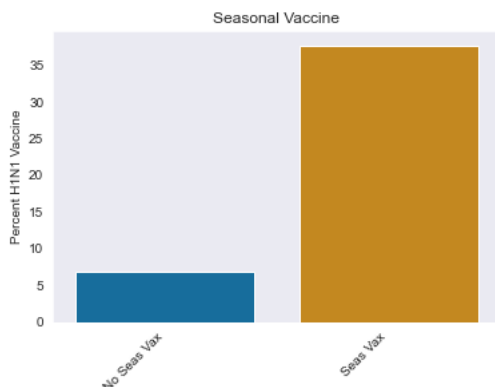
```
0    0.534392
1    0.465608
Name: seasonal_vaccine, dtype: float64
```



```
In [113]: 1 h1n1_rates_house_seas_vax = df.groupby('seasonal_vaccine').mean().h1n1_vaccine
2 print(h1n1_rates_house_seas_vax, ': ', h1n1_rates_house_seas_vax.max() - \
3       h1n1_rates_house_seas_vax.min())
```

```
seasonal_vaccine
0    0.068456
1    0.377724
Name: h1n1_vaccine, dtype: float64 : 0.3092684481726502
```

```
In [115]: 1 # Visualize percent of no seasonal vaccination and seasonal vaccination with H1N1 vaccination
2
3 fig, ax = plt.subplots()
4
5 sns.barplot(x=h1n1_rates_house_seas_vax.index, y=h1n1_rates_house_seas_vax.values*100)
6
7 ax.set_xlabel('')
8 ax.set_ylabel('Percent H1N1 Vaccine')
9 ax.set_title('Seasonal Vaccine')
10 plt.xticks(rotation=45, ha="right")
11
12 ax.set_xticklabels(['No Seas Vax', 'Seas Vax'])
13 plt.savefig('images/doc_recs_vax_perc.png', bbox_inches='tight', dpi=300)
```



What percentage of each question response recieved the H1N1 vaccine?

```

In [13]: 1 # Get percents of H1N1 vaccinated by response by indicator
2
3 # Get percent of vaccinated Lkjsdafllk
4
5 h1n1_rates_house_adlts = df.groupby('household_adults').mean().h1n1_vaccine
6 h1n1_rates_house_child = df.groupby('household_children').mean().h1n1_vaccine
7 h1n1_rates_for_bar_numerical = [h1n1_rates_house_adlts, h1n1_rates_house_child]
8 h1n1_rates_beh_virmedi = df.groupby('behavioral_antiviral_meds').mean().h1n1_vaccine
9 h1n1_rates_beh_avoid = df.groupby('behavioral_avoidance').mean().h1n1_vaccine
10 h1n1_rates_beh_facemask = df.groupby('behavioral_face_mask').mean().h1n1_vaccine
11 h1n1_rates_beh_washhands = df.groupby('behavioral_wash_hands').mean().h1n1_vaccine
12 h1n1_rates_beh_gatherings = df.groupby('behavioral_large_gatherings').mean().h1n1_vaccine
13 h1n1_rates_beh_outside = df.groupby('behavioral_outside_home').mean().h1n1_vaccine
14 h1n1_rates_beh_touchface = df.groupby('behavioral_touch_face').mean().h1n1_vaccine
15 h1n1_rates_docrec = df.groupby('doctor_recc_h1n1').mean().h1n1_vaccine
16 h1n1_rates_chroncond = df.groupby('chronic_med_condition').mean().h1n1_vaccine
17 h1n1_rates_childund6 = df.groupby('child_under_6_months').mean().h1n1_vaccine
18 h1n1_rates_healthworker = df.groupby('health_worker').mean().h1n1_vaccine
19 h1n1_rates_insurance = df.groupby('health_insurance').mean().h1n1_vaccine
20 h1n1_rates_age = df.groupby('age_group').mean().h1n1_vaccine
21 h1n1_rates_education = df.groupby('education').mean().h1n1_vaccine
22 h1n1_rates_race = df.groupby('race').mean().h1n1_vaccine
23 h1n1_rates_sex = df.groupby('sex').mean().h1n1_vaccine
24 h1n1_rates_income = df.groupby('income_poverty').mean().h1n1_vaccine
25 h1n1_rates_marital = df.groupby('marital_status').mean().h1n1_vaccine
26 h1n1_rates_own = df.groupby('rent_or_own').mean().h1n1_vaccine
27 h1n1_rates_employment_status = df.groupby('employment_status').mean().h1n1_vaccine
28 h1n1_rates_geo = df.groupby('hhs_geo_region').mean().h1n1_vaccine.sort_values(ascending=False)
29 h1n1_rates_census = df.groupby('census_msa').mean().h1n1_vaccine
30 h1n1_rates_employment_industry = df.groupby('employment_industry').mean().h1n1_vaccine.sort_values(ascending=False)
31 h1n1_rates_employment_occupation = df.groupby('employment_occupation').mean().h1n1_vaccine.sort_values(ascending=False)
32 h1n1_rates_houseadlt = df.groupby('household_adults').mean().h1n1_vaccine
33 h1n1_rates_housechld = df.groupby('household_children').mean().h1n1_vaccine
34 h1n1_rates_concern = df.groupby('h1n1_concern').mean().h1n1_vaccine
35 h1n1_rates_knowledge = df.groupby('h1n1_knowledge').mean().h1n1_vaccine
36 h1n1_rates_op_effective = df.groupby('opinion_h1n1_vacc_effective').mean().h1n1_vaccine
37 h1n1_rates_op_risk = df.groupby('opinion_h1n1_risk').mean().h1n1_vaccine
38 h1n1_rates_op_sickfromvac = df.groupby('opinion_h1n1_sick_from_vacc').mean().h1n1_vaccine
39
40
41
42 h1n1_rates_for_bars = [h1n1_rates_beh_virmedi, h1n1_rates_beh_avoid,
43                        h1n1_rates_beh_facemask, h1n1_rates_beh_washhands,
44                        h1n1_rates_beh_gatherings, h1n1_rates_beh_outside,
45                        h1n1_rates_beh_touchface, h1n1_rates_docrec,
46                        h1n1_rates_chroncond, h1n1_rates_childund6,
47                        h1n1_rates_healthworker, h1n1_rates_insurance,
48                        h1n1_rates_beh_virmedi, h1n1_rates_beh_avoid,
49                        h1n1_rates_beh_facemask, h1n1_rates_beh_washhands,
50                        h1n1_rates_beh_gatherings, h1n1_rates_beh_outside,
51                        h1n1_rates_beh_touchface, h1n1_rates_docrec,
52                        h1n1_rates_chroncond, h1n1_rates_childund6,
53                        h1n1_rates_healthworker, h1n1_rates_insurance,
54                        h1n1_rates_age, h1n1_rates_education,
55                        h1n1_rates_race, h1n1_rates_sex,
56                        h1n1_rates_income, h1n1_rates_marital,
57                        h1n1_rates_own, h1n1_rates_employment_status,
58                        h1n1_rates_geo, h1n1_rates_census,
59                        h1n1_rates_employment_industry, h1n1_rates_employment_occupation,
60                        h1n1_rates_houseadlt, h1n1_rates_housechld,
61                        h1n1_rates_concern, h1n1_rates_knowledge,
62                        h1n1_rates_op_effective, h1n1_rates_op_risk,
63                        h1n1_rates_op_sickfromvac
64                        ]

```

```
In [14]: 1 # Print percentages by choice and difference between highest and lowest percent by indicator
2
3 for i in h1n1_rates_for_bars:
4     print(i, ': ', i.max() - i.min(), '\n')

4.0    0.176410
5.0    0.404828
Name: h1n1_vaccine, dtype: float64 : 0.35742429292073363


opinion_h1n1_risk
1.0    0.088340
2.0    0.167960
3.0    0.173679
4.0    0.392102
5.0    0.510857
Name: h1n1_vaccine, dtype: float64 : 0.42251705193688244

opinion_h1n1_sick_from_vacc
1.0    0.204601
2.0    0.173184
3.0    0.081081
4.0    0.264274
5.0    0.280293
Name: h1n1_vaccine, dtype: float64 : 0.19921155723624862
```

Get another look at indicators with the largest difference in percentage between lowest and highest percentage choices with bar plot visualizations.

```
In [26]: 1 # Create vizualizations of percentages by choice by indicators
```

```
In [27]: 1 # Visualize proportion of no doc recomend and doc recomend with H1N1 vaccination
2
3 fig, ax = plt.subplots()
4
5 sns.barplot(x=h1n1_rates_docrec.index, y=h1n1_rates_docrec.values*100)
6
7 ax.set_xlabel('')
8 ax.set_ylabel('Percent H1N1 Vaccine')
9 ax.set_title('Doctor Recommendation')
10 plt.xticks(rotation=45, ha="right")
11
12 ax.set_xticklabels(['No Doc Rec', 'Doc Rec'])
13 plt.savefig('images/doc_recs_vax_perc.png', bbox_inches='tight', dpi=300)
```



Who did doctors recommend get the vaccine?

```
In [28]: 1 # Get percents of doctor response by response by indicator
2
3 doctor_recc_rates_age = df.groupby('age_group').mean().doctor_recc_h1n1
4 print(doctor_recc_rates_age, '\n')
5 doctor_recc_rates_education = df.groupby('education').mean().doctor_recc_h1n1
6 print(doctor_recc_rates_education, '\n')
7 doctor_recc_rates_race = df.groupby('race').mean().doctor_recc_h1n1
8 print(doctor_recc_rates_race, '\n')
9 doctor_recc_rates_sex = df.groupby('sex').mean().doctor_recc_h1n1
10 print(doctor_recc_rates_sex, '\n')
11 doctor_recc_rates_income = df.groupby('income_poverty').mean().doctor_recc_h1n1
12 print(doctor_recc_rates_income, '\n')
13 doctor_recc_rates_marital = df.groupby('marital_status').mean().doctor_recc_h1n1
```

```
age_group
18 - 34 Years    0.217536
35 - 44 Years    0.221251
45 - 54 Years    0.203341
55 - 64 Years    0.234421
65+ Years        0.223545
Name: doctor_recc_h1n1, dtype: float64
```

```
education
12 Years          0.202480
< 12 Years        0.200841
College Graduate   0.228896
Some College       0.232847
Name: doctor_recc_h1n1, dtype: float64
```

```
race
Black             0.229102
Hispanic          0.242517
Other or Multiple 0.230717
White             0.216785
Name: doctor_recc_h1n1, dtype: float64
```

```
sex
Female    0.234831
Male      0.198994
Name: doctor_recc_h1n1, dtype: float64
```

```
income_poverty
<= $75,000, Above Poverty    0.216136
> $75,000                    0.235202
Below Poverty                 0.235459
Name: doctor_recc_h1n1, dtype: float64
```

Nothing jumps out as an imbalance in percent of doctor recommendations, but all percentages are low. What was the overall percentage of those that responded when they received a doctor recommendation?




```
In [29]: 1 # Look at the percent of respondents who recieved a doctor recommendation to get the vaccine
2
3 # Get percents
4 doc_recc_prop = df.doctor_recc_h1n1.value_counts(normalize=True)
5 print(doc_recc_prop)
6
7
8 # Chart
9 fig, ax = plt.subplots()
10
11 sns.barplot(x=doc_recc_prop.index, y=doc_recc_prop.values*100)
12
13 ax.set_xlabel('')
14 ax.set_ylabel('Percent Doc Recommendation')
15 ax.set_title("Doctor Recommendation for Vaccine")
16
17 plt.xticks(rotation=45, ha="right")
18
19 ax.set_xticklabels(['No Rec', 'Rec'])
20
21 plt.savefig('images/perc_doc_recs.png', bbox_inches='tight', dpi=300)
```


0.0 0.779688

1.0 0.220312

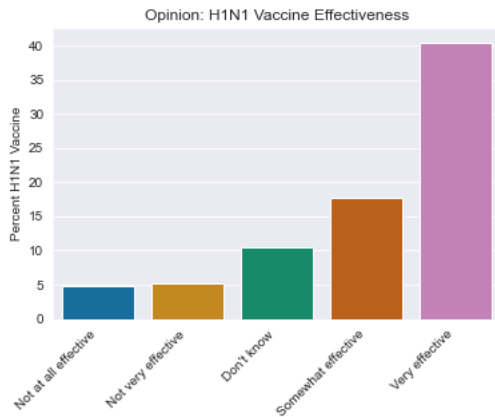
Name: doctor\_recc\_h1n1, dtype: float64



```
In [30]: 1 # Visualize proportion of level of respons opinion about risk of getting sick with H1N1 without vaccine with H1N1 vaccine
2
3 fig, ax = plt.subplots()
4
5 sns.barplot(x=h1n1_rates_op_risk.index, y=h1n1_rates_op_risk.values*100)
6
7 ax.set_xlabel('')
8 ax.set_ylabel('Percent H1N1 Vaccine')
9 ax.set_title('Opinion: Risk of Sick Without Vaccine')
10 plt.xticks(rotation=45, ha="right")
11
12 ax.set_xticklabels(['Very low', 'Somewhat low', "Don't know", 'Somewhat high', 'Very high'])
13
14 plt.savefig('images/op_sick_without_vax.png', bbox_inches='tight', dpi=300)
```



```
In [31]: 1 # Visualize proportion of level of respons opinion about seasonal flu vaccine effectiveness with H1N1 vaccination
2
3 fig, ax = plt.subplots()
4
5 sns.barplot(x=h1n1_rates_op_effective.index, y=h1n1_rates_op_effective.values*100)
6
7 ax.set_xlabel('')
8 ax.set_ylabel('Percent H1N1 Vaccine')
9 ax.set_title('Opinion: H1N1 Vaccine Effectiveness')
10 plt.xticks(rotation=45, ha="right")
11
12 ax.set_xticklabels(['Not at all effective', 'Not very effective', "Don't know", 'Somewhat effective', 'Very effective'])
13
14 plt.savefig('images/op_vax_effective.png', bbox_inches='tight', dpi=300)
```

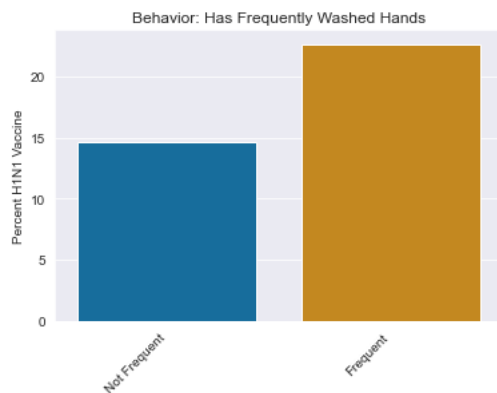


```
In [32]: 1 # Visualize proportion of level of respons opinion about seasonal flu vaccine effectiveness with H1N1 vaccination
2
3 fig, ax = plt.subplots()
4
5 sns.barplot(x=h1n1_rates_beh_facemask.index, y=h1n1_rates_beh_facemask.values*100)
6
7 ax.set_xlabel('')
8 ax.set_ylabel('Percent H1N1 Vaccine')
9 ax.set_title('Behavior: Purchased Facemask')
10 plt.xticks(rotation=45, ha="right")
11
12 ax.set_xticklabels(['No Facemask', 'Purchased Facemask'])
13
14 plt.savefig('images/beh_facemask.png', bbox_inches='tight', dpi=300)
```



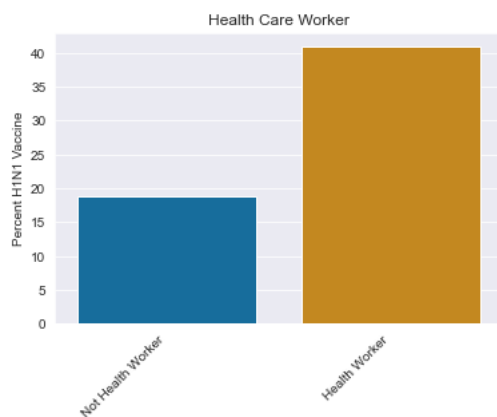
```
In [33]: 1 # Visualize proportion of level of respons opinion about seasonal flu vaccine effectiveness with H1N1 vaccination
2
3 fig, ax = plt.subplots()
4
5 sns.barplot(x=h1n1_rates_beh_washhands.index, y=h1n1_rates_beh_washhands.values*100)
6
7 ax.set_xlabel('')
8 ax.set_ylabel('Percent H1N1 Vaccine')
9 ax.set_title('Behavior: Has Frequently Washed Hands')
10 plt.xticks(rotation=45, ha="right")
11
12 ax.set_xticklabels(['Not Frequent', 'Frequent'])
```

Out[33]: [Text(0, 0, 'Not Frequent'), Text(1, 0, 'Frequent')]



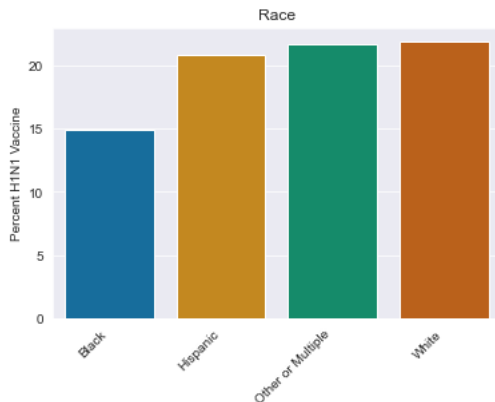
```
In [34]: 1 # Visualize proportion of not health care worker and health care worker with H1N1 vaccination
2
3 fig, ax = plt.subplots()
4
5 sns.barplot(x=h1n1_rates_healthworker.index, y=h1n1_rates_healthworker.values*100)
6
7 ax.set_xlabel('')
8 ax.set_ylabel('Percent H1N1 Vaccine')
9 ax.set_title('Health Care Worker')
10 plt.xticks(rotation=45, ha="right")
11
12 ax.set_xticklabels(['Not Health Worker', 'Health Worker'])
```

Out[34]: [Text(0, 0, 'Not Health Worker'), Text(1, 0, 'Health Worker')]



```
In [35]: 1 # Visualize proportion of not health care worker and health care worker with H1N1 vaccination
2
3 fig, ax = plt.subplots()
4
5 sns.barplot(x=h1n1_rates_race.index, y=h1n1_rates_race.values*100)
6
7 ax.set_xlabel('')
8 ax.set_ylabel('Percent H1N1 Vaccine')
9 ax.set_title('Race')
10 plt.xticks(rotation=45, ha="right")
```

```
Out[35]: (array([0, 1, 2, 3]),
 [Text(0, 0, 'Black'),
  Text(1, 0, 'Hispanic'),
  Text(2, 0, 'Other or Multiple'),
  Text(3, 0, 'White')])
```



Let's have some fun, as a reminder that human behavior is a difficult thing to predict. Are there respondents that are surprising to find did not receive the H1N1 vaccine? Let's see if anyone who gave the most vaccine positive responses to the questions likeliness to get sick from vaccine, the vaccine is effective, and received a doctor's recommendation and still did not get an H1N1 vaccine.

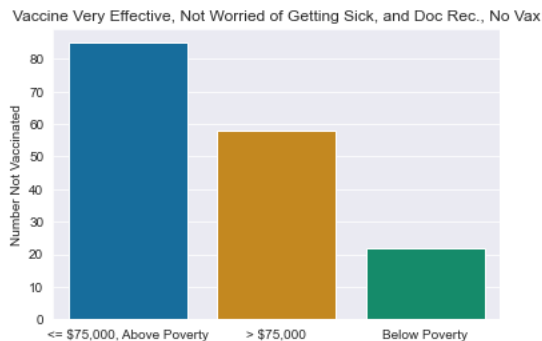
```
In [37]: 1 df_y_no_vax = df[(df['h1n1_vaccine'] == 0) &
2                       (df['opinion_h1n1_sick_from_vacc'] == 1) &
3                       (df['opinion_h1n1_vacc_effective'] == 5) &
4                       (df['doctor_recc_h1n1'] == 1)]
```

```
In [38]: 1 y_no_vax_norm = df_y_no_vax.income_poverty.value_counts(normalize=True)
2 print(y_no_vax_norm, '\n')
3 y_no_vax = df_y_no_vax.income_poverty.value_counts()
4 print(y_no_vax)
```

```
<= $75,000, Above Poverty    0.515152
> $75,000                    0.351515
Below Poverty                0.133333
Name: income_poverty, dtype: float64
```

```
<= $75,000, Above Poverty    85
> $75,000                    58
Below Poverty                22
Name: income_poverty, dtype: int64
```

```
In [39]: 1 # Chart
2 fig, ax = plt.subplots()
3
4 sns.barplot(x=y_no_vax.index, y=y_no_vax.values)
5
6 ax.set_xlabel('')
7 ax.set_ylabel('Number Not Vaccinated')
8 ax.set_title("Vaccine Very Effective, Not Worried of Getting Sick, and Doc Rec., No Vax")
9
10 ax.set_xticklabels(['<= $75,000, Above Poverty', '> $75,000', 'Below Poverty',])
11
12 plt.savefig('images/###.png', bbox_inches='tight', dpi=300)
```



```
In [ ]: 1 # Proportion of those who *%*%*^% who got h1n1 vaccine
2
3 likely_vax.groupby('income_poverty').mean().h1n1_vaccine
```

## Modeling

1 Having explored the data some, build first models; a dummy model and first simple model.

```
In [41]: 1 X = df.drop(['respondent_id', 'h1n1_vaccine'], axis=1)
2 y = df.h1n1_vaccine
3
4 X_train, X_test, y_train, y_test = train_test_split(X,y,random_state=42)
```

```
In [42]: 1 print(X_train.shape)
2 print(y_train.shape)
```

```
(20030, 36)
(20030,)
```

Set up transformer pipelines and a cross validation class for modeling.

```
In [43]: 1 # Numeric Pipeline, impute missing data and scale
2
3 num_pipe = Pipeline([
4     ('num_impute', SimpleImputer(strategy='mean')),
5     ('scaler', StandardScaler())
6 ])
7
8 # Categorical Pipeline, impute missing data and encode categoricals
9
10 cat_pipe = Pipeline([
11     ('cat_impute', SimpleImputer(strategy='most_frequent')),
12     ('one', OneHotEncoder(sparse=False, handle_unknown='ignore'))
13 ])
14
15 # Create column transformer to use in model pipelines
16
17 CT = ColumnTransformer([
18     ('num_trans', num_pipe, selector(dtype_include=np.number)),
19     ('cat_trans', cat_pipe, selector(dtype_include=object)),
20 ], remainder='passthrough')
```

```

In [183]: 1 class ModelWithCV():
2         '''Structure to save the model and more easily see its crossvalidation'''
3
4         def __init__(self, model, model_name, X, y, cv_now=True):
5             self.model = model
6             self.name = model_name
7             self.X = X
8             self.y = y
9             # For CV results
10            self.cv_results = None
11            self.cv_mean = None
12            self.cv_median = None
13            self.cv_std = None
14            #
15            if cv_now:
16                self.cross_validate()
17
18            def cross_validate(self, X=None, y=None, kfold=10):
19                '''
20                Perform cross-validation and return results.
21
22                Args:
23                X:
24                    Optional; Training data to perform CV on. Otherwise use X from object
25                y:
26                    Optional; Training data to perform CV on. Otherwise use y from object
27                kfold:
28                    Optional; Number of folds for CV (default is 10)
29                '''
30
31                cv_X = X if X else self.X
32                cv_y = y if y else self.y
33
34                self.cv_results = cross_val_score(self.model, cv_X, cv_y, scoring='recall', cv=kfold)
35                self.cv_mean = np.mean(self.cv_results)
36                self.cv_median = np.median(self.cv_results)
37                self.cv_std = np.std(self.cv_results)
38
39            def print_cv_summary(self):
40                cv_summary = (
41                    f'''CV Results for `{self.name}` model:
42                    {self.cv_mean:.5f} ± {self.cv_std:.5f} recall accuracy
43                    '''
44                )
45                print(cv_summary)
46
47            def plot_cv(self, ax):
48                '''
49                Plot the cross-validation values using the array of results and given
50                Axis for plotting.
51                '''
52
53                ax.set_title(f'CV Results for `{self.name}` Model')
54                # Thinner violinplot with higher bw
55                sns.violinplot(y=self.cv_results, ax=ax, bw=.4)
56                sns.swarmplot(
57                    y=self.cv_results,
58                    color='orange',
59                    size=10,
60                    alpha=0.8,
61                    ax=ax
62                )
63
64                return ax

```

- 1 There is a class imbalance for the target, use SMOTE to create synthetic data to balance the classes.
- 2 Use ModelWithCV class to cross validate.

```

In [185]: 1 # Create Dummy/Baseline
2
3 dummy_smote_model = ImPipeline([
4     ('ct', CT),
5     ('sm', SMOTE(random_state=42)),
6     ('dummy', DummyClassifier(strategy='most_frequent', random_state=42))
7 ])
8
9 # Use the class with our dummy pipe
10
11 dummy_smoted_model_pipe = ModelWithCV(dummy_smote_model, model_name='dummy_smote', X=X_train, y=y_train)

```

```
In [186]: 1 dummy_smoted_model_pipe.cv_mean
```

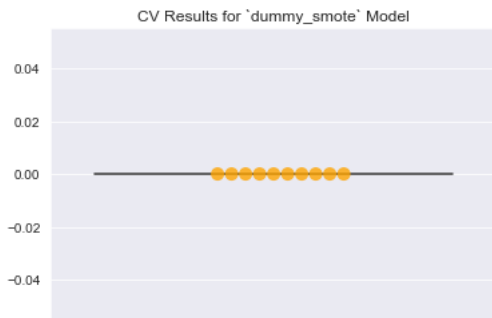
```
Out[186]: 0.0
```

```
In [187]: 1 dummy_smoted_model_pipe.print_cv_summary()
```

```
CV Results for `dummy_smote` model:
0.00000 ± 0.00000 recall accuracy
```

```
In [188]: 1 sns.set_style("darkgrid")
2 fig,ax = plt.subplots()
3
4 dummy_smoted_model_pipe.plot_cv(ax=ax)
```

```
Out[188]: <AxesSubplot:title={'center':'CV Results for `dummy_smote` Model'}>
```



Our dummy model should have given us a .50 after creating synthetic data, however it did not. This is something to look into fixing. Going with what it produced without synthetic data, it is expected to have a recall score of 0 as it only predicted the most frequent outcome of not getting the vaccine.

The first simple logistic regression model will be sent all data and use SMOTE to balance the classes.

```
In [191]: 1 # Use the class with Logreg pipe
2
3 logreg_fsm = ImPipeline([
4     ('ct', CT),
5     ('sm', SMOTE(random_state=42)),
6     ('logreg_fsm', LogisticRegression(random_state=42, max_iter=1000))
7 ])
```


```
In [192]: 1 fsm_model_pipe = ModelWithCV(logreg_fsm, model_name='fsm', X=X_train, y=y_train)
```

```
In [193]: 1 fsm_model_pipe.print_cv_summary()
```

```
CV Results for `fsm` model:
0.79422 ± 0.02376 recall accuracy
```

```
In [194]: 1 fig,ax = plt.subplots()
2
3 fsm_model_pipe.plot_cv(ax=ax)
```

```
Out[194]: <AxesSubplot:title={'center':'CV Results for `fsm` Model'}>
```



Make adjustments to the features, drop those with high percentage of nulls, and sum behavioral group to reduce features and create stronger indicator. To reduce features, combine the behavioral questions into a behavioral sum category.

```

In [195]: 1 # Create Behavioral Sum
2
3 sns.set_style("darkgrid")
4
5 #Proportional Behavioral Sum
6 df['beh_sum'] = df.behavioral_antiviral_meds + df.behavioral_avoidance + \
7               df.behavioral_face_mask + df.behavioral_wash_hands + \
8               df.behavioral_large_gatherings + df.behavioral_outside_home + \
9               df.behavioral_touch_face
10
11 beh_sum = df.groupby('beh_sum').mean().h1n1_vaccine
12 print(beh_sum)
13 print(beh_sum.max() - beh_sum.min())
14
15 fig, ax = plt.subplots()
16
17 sns.barplot(x=beh_sum.index, y=beh_sum.values*100)
18
19 ax.set_xlabel('Behavioral Sum')
20 ax.set_ylabel('Percent H1N1 Vaccine')
21 ax.set_title('Total Behavior Score')
22 plt.xticks(rotation=45, ha="right")
23
24 plt.savefig('images/beh_totals.png', bbox_inches='tight', dpi=300)

```

```

beh_sum
0.0    0.120429
1.0    0.162809
2.0    0.207301
3.0    0.227935
4.0    0.227830
5.0    0.230254
6.0    0.284333
7.0    0.333333
Name: h1n1_vaccine, dtype: float64
0.2129042743377214

```

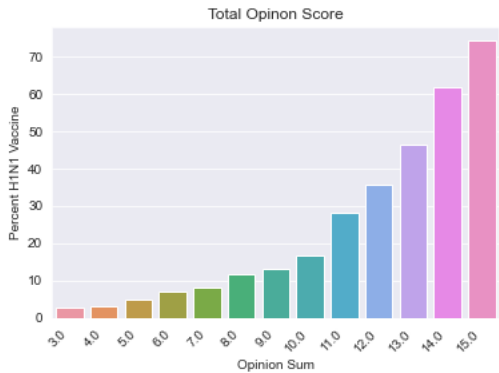




In [196]:

```
1 # Ceate Opinion Sum
2
3 #Proportional Opinion Sum
4 df['op_sum'] = df.opinion_h1n1_vacc_effective + df.opinion_h1n1_risk + \
5               (6 - df.opinion_h1n1_sick_from_vacc)
6
7 op_sum = df.groupby('op_sum').mean().h1n1_vaccine
8 print(op_sum)
9 print(op_sum.max() - op_sum.min())
10
11 fig, ax = plt.subplots()
12
13 sns.barplot(x=op_sum.index, y=op_sum.values*100)
14
15 ax.set_xlabel('Opinion Sum')
16 ax.set_ylabel('Percent H1N1 Vaccine')
17 ax.set_title('Total Opinon Score')
18 plt.xticks(rotation=45, ha="right")
19
20 plt.savefig('images/op_totals.png', bbox_inches='tight', dpi=300)
```

```
op_sum
3.0    0.026316
4.0    0.032051
5.0    0.049470
6.0    0.070288
7.0    0.081171
8.0    0.115534
9.0    0.132679
10.0   0.165134
11.0   0.282563
12.0   0.357179
13.0   0.463612
14.0   0.617964
15.0   0.742308
Name: h1n1_vaccine, dtype: float64
0.7159919028340082
```



In [197]:

```
1 # Create dataframe with new features
2
3 df_fm = df.copy(['h1n1_concern', 'h1n1_knowledge', 'beh_sum', \
4                  'doctor_recc_h1n1', 'chronic_med_condition', 'child_under_6_months', \
5                  'health_worker', 'opinion_h1n1_vacc_effective', 'opinion_h1n1_risk', \
6                  'opinion_h1n1_sick_from_vacc', 'age_group', 'education', 'race', 'sex', \
7                  'income_poverty', 'marital_status', 'rent_or_own', 'employment_status', \
8                  'hhs_geo_region', 'household_adults', 'household_children', 'seasonal_vaccine', \
9                  'h1n1_vaccine'])
```

In [199]:

```
1 df_fm.head()
```

Out[199]:

	respondent_id	h1n1_concern	h1n1_knowledge	behavioral_antiviral_meds	behavioral_avoidance	behavioral_face_mask	behavioral_wash_hands	behavior
0	0	1.0	0.0	0.0	0.0	0.0	0.0	
1	1	3.0	2.0	0.0	1.0	0.0	1.0	
2	2	1.0	1.0	0.0	1.0	0.0	0.0	
3	3	1.0	1.0	0.0	1.0	0.0	1.0	
4	4	2.0	1.0	0.0	1.0	0.0	1.0	

```
In [66]: 1 # Create new train test split with new dataframe
2
3 X = df_fm.drop(['respondent_id', 'h1n1_vaccine'], axis=1)
4 y = df_fm.h1n1_vaccine
5
6 X_clean_train, X_clean_test, y_clean_train, y_clean_test = train_test_split(X,y,random_state=42)
```

The next model will be the same model type and parameters with the adjusted data.

```
In [200]: 1 fsm_clean_model_pipe = ModelWithCV(logreg_fsm, model_name='fsm_clean', X=X_clean_train, y=y_clean_train)
```


```
In [201]: 1 fsm_clean_model_pipe.cv_mean
```

```
Out[201]: 0.7949229494614747
```

```
In [ ]: 1 fsm_clean_model_pipe.
```

```
In [202]: 1 fig,ax = plt.subplots()
2
3 fsm_model_pipe.plot_cv(ax=ax)
```

```
Out[202]: <AxesSubplot:title={'center':'CV Results for `fsm` Model'}>
```



1 This simple model with the adjusted features produced almost exactly the same results, but with less features will help run future models faster.

Use grid search to find the best parameters for a logistic regression model.

```
In [204]: 1 # Logistic Regression
2
3 lr = LogisticRegression(random_state=42, max_iter=1000)
4
5 lr_model_pipe = ImPipeline([
6     ('ct', CT),
7     ('sm', SMOTE(random_state=42)),
8     ('lr', lr)
9 ])
10
11 lr_params = {
12     'ct_num_trans_num_impute_strategy' : ['mean', 'median'],
13     'lr_penalty' : ['l1', 'l2', 'elasticnet'],
14     'lr_C' : [100, 10, 1.0, 0.1, 0.01],
15     'lr_solver' : ['lbfgs', 'liblinear', 'saga'],
16 }
```

```
In [205]: 1 gs_lr = GridSearchCV(estimator=lr_model_pipe, param_grid=lr_params, scoring='recall', cv=10, verbose=2)
2
3 gs_lr.fit(X_clean_train, y_clean_train)
```

```
Out[205]: GridSearchCV(cv=10,
                      estimator=Pipeline(steps=[('ct',
                                                  ColumnTransformer(remainder='passthrough',
                                                                      transformers=[('num_trans',
                                                                      Pipeline(steps=[('num_impute',
                                                                      SimpleImputer()),
                                                                      ('scaler',
                                                                      StandardScaler()))],
                                                                      <sklearn.compose._column_transformer.make_column
                                                                      _selector object at 0x0000028A1471A790>),
                                                  ('cat_trans',
                                                  Pipeline(steps=[('cat_impute',
                                                                      SimpleImputer(strategy='...
                                                                      <sklearn.compose._column_transformer.make_column
                                                                      _selector object at 0x0000028A1471A850>)))]),
                      ('sm', SMOTE(random_state=42)),
                      ('lr',
                      LogisticRegression(max_iter=1000,
                                          random_state=42))]),
                      param_grid={'ct__num_trans__num_impute__strategy': ['mean',
```

```
In [146]: 1 gs_lr.best_params_
```

```
Out[146]: {'ct__num_trans__num_impute__strategy': 'mean',
           'lr__C': 0.01,
           'lr__penalty': 'l1',
           'lr__solver': 'liblinear'}
```

```
In [147]: 1 gs_lr.best_score_
```

```
Out[147]: 0.8076128141397405
```

```
1 Use grid search to find the best parameters for a random forest classifier model.
```

```
In [150]: 1 # Random Forest Classifier
2
3 rfc = RandomForestClassifier(random_state=42)
4
5 rfc_model_pipe = ImPipeline([
6     ('ct', CT),
7     ('sm', SMOTE(random_state=42)),
8     ('rfc', rfc)])
9
10 rfc_params = {
11     'ct__num_trans__num_impute__strategy': ['mean', 'median'],
12     'rfc__criterion': ['gini', 'entropy', 'log_loss'],
13     'rfc__max_depth': [9, 11, 13, 15],
14     'rfc__min_samples_leaf': [1, 5, 15],
15 }
```

```
In [151]: 1 gs_rfc = GridSearchCV(estimator=rfc_model_pipe, param_grid=rfc_params, scoring='recall', cv=10, verbose=2)
2
3 gs_rfc.fit(X_clean_train, y_clean_train)
```

```
Out[151]: GridSearchCV(cv=10,
                      estimator=Pipeline(steps=[('ct',
                                                  ColumnTransformer(remainder='passthrough',
                                                                      transformers=[('num_trans',
                                                                      Pipeline(steps=[('num_impute',
                                                                      SimpleImputer()),
                                                                      ('scaler',
                                                                      StandardScaler()))],
                                                                      <sklearn.compose._column_transformer.make_column
                                                                      _selector object at 0x0000028A1471A790>),
                                                  ('cat_trans',
                                                  Pipeline(steps=[('cat_impute',
                                                                      SimpleImputer(strategy='...
                                                                      <sklearn.compose._column_transformer.make_column
                                                                      _selector object at 0x0000028A1471A850>)))]),
                      ('sm', SMOTE(random_state=42)),
                      ('rfc',
                      RandomForestClassifier(random_state=42))]),
                      param_grid={'ct__num_trans__num_impute__strategy': ['mean',
```

In [152]: 1 gs\_rfc.best\_params\_

```
Out[152]: {'ct__num_trans__num_impute__strategy': 'mean',
'rfc__criterion': 'entropy',
'rfc__max_depth': 11,
'rfc__min_samples_leaf': 15}
```

In [153]: 1 gs\_rfc.best\_score\_

```
Out[153]: 0.62602209334438
```

1 Because the best parameters had a parameter, min\_sample\_leaf, at the far end of our options, run another model with the best parameters but increasing the options for min\_sample\_leaf.

In [ ]: 1

```
In [174]: 1 rfc_params_adj_leafs = {
2         'ct__num_trans__num_impute__strategy' : ['mean'],
3         'rfc__criterion' : ['entropy'],
4         'rfc__max_depth' : [11],
5         'rfc__min_samples_leaf' : [15, 17, 19, 21, 23]
6     }
```

```
In [175]: 1 gs_rfc = GridSearchCV(estimator=rfc_model_pipe, param_grid=rfc_params_adj_leafs, scoring='recall', cv=10, verbose=2)
2
3 gs_rfc.fit(X_clean_train, y_clean_train)
```

```
transformers=[('num_trans',
               Pipeline(steps=[('num_impute',
                                SimpleImputer()),
                                ('scaler',
                                 StandardScaler())])),
              <sklearn.compose._column_transformer.make_column
_selector object at 0x0000028A1471A790>),
              ('cat_trans',
               Pipeline(steps=[('cat_impute',
                                SimpleImputer(strategy='...
                                sparse=False))])),
              <sklearn.compose._column_transformer.make_column
_selector object at 0x0000028A1471A850>))],
              ('sm', SMOTE(random_state=42)),
              ('rfc',
               RandomForestClassifier(random_state=42))]),
              param_grid={'ct__num_trans__num_impute__strategy': ['mean'],
                           'rfc__criterion': ['entropy'], 'rfc__max_depth': [11],
                           'rfc__min_samples_leaf': [15, 17, 19, 21, 23]},
              scoring='recall', verbose=2)
```

In [179]: 1 gs\_rfc.best\_params\_

```
Out[179]: {'ct__num_trans__num_impute__strategy': 'mean',
'rfc__criterion': 'entropy',
'rfc__max_depth': 11,
'rfc__min_samples_leaf': 21}
```

In [176]: 1 gs\_rfc.best\_score\_

```
Out[176]: 0.6295476387738195
```

In [ ]: 1

1 Use grid search to find the best parameters for a k-nearest neighbors classifier model.

```

In [180]: 1 # K-Nearest Neighbors Classifier
2
3 knn = KNeighborsClassifier()
4
5 knn_model_pipe = ImPipeline([
6     ('ct', CT),
7     ('sm', SMOTE(random_state=42)),
8     ('knn', knn)])
9
10 knn_params = {
11     'ct_num_trans_num_impute_strategy' : ['mean', 'median'],
12     'knn_n_neighbors' : [1, 3, 5, 7, 9, 11],
13     'knn_weights' : ['uniform', 'distance'],
14     'knn_leaf_size' : [5, 10, 20, 30, 40],
15     'knn_metric' : ['minkowski', 'cityblock']
16 }

In [181]: 1 gs_knn = GridSearchCV(estimator=knn_model_pipe, param_grid=knn_params, scoring='recall', cv=10, verbose=2)
2
3 gs_knn.fit(X_clean_train, y_clean_train)

```

```

In [ ]: 1 gs_rfc.best_score_

```


## Evlauation

The logistic regression model with a numerical impute strategy of 'mean, C: .01, penalty: l1, and solver: liblinear, was our best performing model on trained data and cross validation. Now compare the test data.

```

In [210]: 1 sns.set_style("dark")
2 gs_lr.best_estimator_.fit(X_clean_train, y_clean_train)
3 y_pred = gs_lr.best_estimator_.predict(X_clean_test)
4 plot_confusion_matrix(estimator=gs_lr.best_estimator_, X=X_clean_test, y_true=y_clean_test)
5
6 plt.savefig('images/cm_log_reg.png', bbox_inches='tight', dpi=300)

```



```

In [211]: 1 recall_score(y_clean_test, y_pred)

```

Out[211]: 0.8038108680310515

```

In [212]: 1 precision_score(y_clean_test, y_pred)

```

Out[212]: 0.5135256988277728